

CS 6V81-05: System Security and Malicious Code Analysis

Buffer Overflow, Integer and Heap Overflow

Zhiqiang Lin

Department of Computer Science
University of Texas at Dallas

April 11th, 2012

Outline

- 1 Background
- 2 Buffer Overflow
- 3 Integer and Heap Overflow
- 4 Summary

- 1 Background
- 2 Buffer Overflow
- 3 Integer and Heap Overflow
- 4 Summary


IA32 Stack

Region of memory
managed with stack
discipline

Grows toward lower
addresses

Register `%esp` indicates
lowest stack address

address of top element



IA32 Stack Pushing


Pushing

```
pushl src
```

Fetch operand at Src

Decrement %esp by 4

Write operand at
address given by %esp



IA32 Stack Popping


Popping

```
popl dest
```

Read operand at
address given by %esp

Increment %esp by 4



Write to Dest



Procedure Call Example

```
804854e: e8 3d 06 00 00  call    8048b90 <main>
8048553: 50              pushl   %eax
```

call 8048b90



%esp: 0x108

%esp: 0x108

%eip: 0x804854e



%eip: 0x804854e

%eip is program counter

Procedure Call Example

```
804854e: e8 3d 06 00 00    call    8048b90 <main>
8048553: 50               pushl   %eax
```

```
call    8048b90
```



%esp	0x108
-------------	--------------

<code>%esp</code>	0x104
-------------------	-------


```
%eip 0x804854e
```

```
%eip 0x804854e
```

%eip is program counter

```
804854e: e8 3d 06 00 00    call    8048b90 <main>
8048553: 50                pushl   %eax
```

```
call    8048b90
```




%eip is program counter

Procedure Call Example


```
804854e: e8 3d 06 00 00  call    8048b90 <main>
8048553: 50              pushl   %eax
```

call 8048b90



`%esp` `0x108`

`%eip` `0x804854e`



`%esp` `0x104`

`%eip` `0x8048b90`


`%eip` is program counter

Procedure Return Example

8048591: c3


ret

ret



%esp 0x104

%eip 0x8048591



%esp 0x108

%eip 0x8048553

%eip is program counter

Call Chain Example

CodeStructure


```
yoo (...)  
{  
  .  
  .  
  who ();  
  .  
  .  
}
```

```
who (...)  
{  
  . . .  
  amI ();  
  . . .  
  amI ();  
  . . .  
}
```

```
amI (...)  
{  
  .  
  .  
  amI ();  
  .  
  .  
}
```

Procedure amI recursive

Call Chain



Call Chain Example

Contents

- Local variables
- Return information
- Temporary space

Management

- Space allocated when enter procedure

- “Set-up” code


- Deallocated when return

- “Finish” code


Pointers

- Stack pointer **%esp** indicates stack top

- Frame pointer **%ebp** indicates start of current frame




Stack Operation




Call Chain


yoo




Stack Operation





Call Chain




Stack Operation




Call Chain




Stack Operation





Call Chain




Stack Operation





Call Chain




Stack Operation





Call Chain




Stack Operation




Call Chain




Stack Operation




Call Chain




Stack Operation




Call Chain



Stack Operation




Call Chain



Stack Operation



Call Chain



IA32/Linux Stack Frame

Current Stack Frame (“Top” to Bottom)

Parameters for function about to call

“Argument build”

Local variables

If can't keep in registers

Saved register context

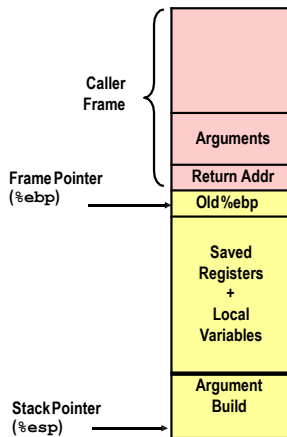
Old frame pointer

Caller Stack Frame

Return address

Pushed by **call** instruction

Arguments for this call



IA32 Linux Memory Layout

Stack

Runtime stack (8MB limit)

E. g., local variables

Heap

Dynamically allocated storage

When call malloc(), calloc(),
new()

Data


Statically allocated data

E.g., arrays & strings declared
in code

Text

Executable machine
instructions

Read-only



Upper 2 hex digits
= 8 bits of address

Memory Allocation Example


```
char big_array[1<<24]; /* 16 MB */
char huge_array[1<<28]; /* 256 MB*/

int beyond;
char *p1, *p2, *p3, *p4;

int useless() { return 0; }

int main()
{
    p1 = malloc(1 <<28); /* 256 MB */
    p2 = malloc(1 << 8); /* 256 B  */
    p3 = malloc(1 <<28); /* 256 MB */
    p4 = malloc(1 << 8); /* 256 B  */
    /* Some print statements ... */
}
```

Where does everything go?




IA32 Example Addresses

address range $\sim 2^{32}$

<code>\$esp</code>	<code>0xfffffbcd0</code>
<code>p3</code>	<code>0x65586008</code>
<code>p1</code>	<code>0x55585008</code>
<code>p4</code>	<code>0x1904a110</code>
<code>p2</code>	<code>0x1904a008</code>
<code>&p2</code>	<code>0x18049760</code>
<code>&beyond</code>	<code>0x08049744</code>
<code>big_array</code>	<code>0x18049780</code>
<code>huge_array</code>	<code>0x08049760</code>
<code>main()</code>	
<code>useless()</code>	<code>0x08049744</code>
<code>final malloc()</code>	<code>0x006be166</code>

`malloc()` is dynamically linked
address determined at runtime



Outline

- 1 Background
- 2 **Buffer Overflow**
- 3 Integer and Heap Overflow
- 4 Summary

Internet Worm and IM War


November, 1988

Internet Worm attacks thousands of Internet hosts.
How did it happen?

July, 1999

Microsoft launches MSN Messenger (instant messaging system).

Messenger clients can access popular AOL Instant Messaging Service (AIM) servers



Internet Worm and IM War (cont.)

August 1999

Mysteriously, Messenger clients can no longer access AIM servers.

Microsoft and AOL begin the IM war:

- AOL changes server to disallow Messenger clients

- Microsoft makes changes to clients to defeat AOL changes.

- At least 13 such skirmishes.

How did it happen?

The Internet Worm and AOL/Microsoft War were both based on stack buffer overflow exploits!

- many library functions do not check argument sizes.

- allows target buffers to overflow.

String Library Code

Implementation of Unix function gets()

```
/* Get string from stdin */
char *gets(char *dest)
{
    int c = getchar();
    char *p = dest;
    while (c != EOF && c != '\n') {
        *p++ = c;
        c = getchar();
    }
    *p = '\0';
    return dest;
}
```

No way to specify limit on number of characters to read

Similar problems with other library functions

strcpy, strcat: Copy strings of arbitrary length

scanf, fscanf, sscanf, when given %s conversion specification

Vulnerable Buffer Code

```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

```
void call_echo() {
    echo();
}
```

```
unix>./bufdemo
Type a string:1234567
1234567
```

```
unix>./bufdemo
Type a string:12345678
Segmentation Fault
```

```
unix>./bufdemo
Type a string:123456789ABC
Segmentation Fault
```

Buffer Overflow Disassembly

echo

```

80485c5: 55          push    %ebp
80485c6: 89 e5       mov     %esp,%ebp
80485c8: 53          push    %ebx
80485c9: 83 ec 14    sub     $0x14,%esp
80485cc: 8d 5d f8    lea     0xffffffff8(%ebp),%ebx
80485cf: 89 1c 24    mov     %ebx, (%esp)
80485d2: e8 9e ff ff call     8048575 <gets>
80485d7: 89 1c 24    mov     %ebx, (%esp)
80485da: e8 05 fe ff call     80483e4 <puts@plt>
80485df: 83 c4 14    add     $0x14,%esp
80485e2: 5b          pop     %ebx
80485e3: 5d          pop     %ebp
80485e4: c3          ret

```

call echo


```

80485eb: e8 d5 ff ff call     80485c5 <echo>
80485f0: c9          leave
80485f1: c3          ret

```

Buffer Overflow Stack

Before call to gets



```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```


```
echo:
pushl %ebp # Save %ebp on stack
movl %esp, %ebp
pushl %ebx # Save %ebx
subl $20, %esp # Allocate stack space
leal -8(%ebp), %ebx # Compute buf as %ebp-8
movl %ebx, (%esp) # Push buf on stack
call gets # Call gets
...
```

Buffer Overflow Stack Example

```

unix> gdb bufdemo
(gdb) break echo
Breakpoint 1 at 0x80485c9
(gdb) run
Breakpoint 1, 0x80485c9 in echo ()
(gdb) print /x $ebp
$1 = 0xffffd678
(gdb) print /x *(unsigned *)$ebp
$2 = 0xffffd688
(gdb) print /x *((unsigned *)$ebp + 1)
$3 = 0x80485f0
  
```


Before call to gets



```

80485eb: e8 d5 ff ff ff
80485f0: c9
  
```

Before call to gets

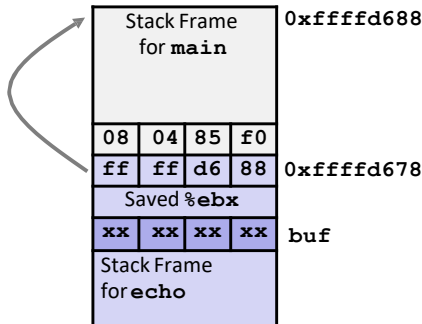


```


call 80485c5 <echo>
leave
  
```

Buffer Overflow Stack Example #1

Before call to gets



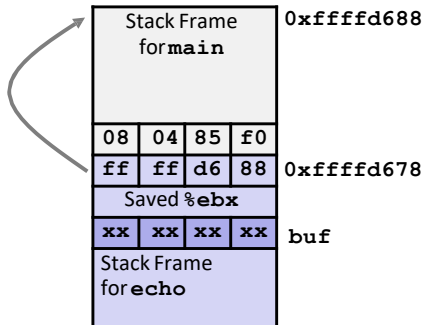
Input 1234567




**Overflow buf, and corrupt `%ebx`,
but no problem**

Buffer Overflow Stack Example #2

Before call to gets



Input 12345678



Base pointer corrupted

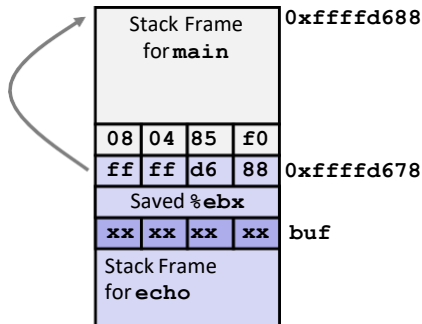
```

. . .
80485eb:  e8 d5 ff ff ff    call    80485c5 <echo>
80485f0:  c9                leave   # Set %ebp to corrupted value
80485f1:  c3                ret


```

Buffer Overflow Stack Example #3

Before call to gets



Input 123456789!"#



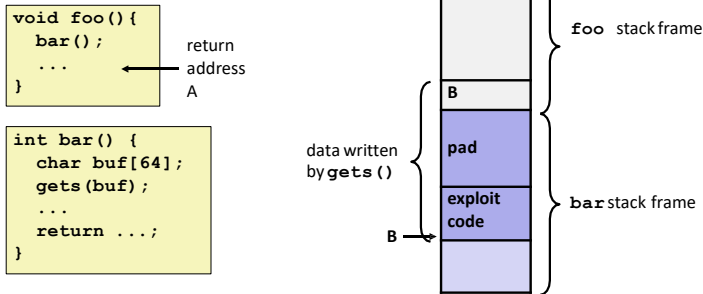
Return address corrupted

```

80485eb:  e8 d5 ff ff ff    call    80485c5 <echo>
80485f0:  c9                leave   # Desired return point

```


Buffer Overflow Stack Example #3



Input string contains byte representation of executable code

Overwrite return address A with address of buffer B

When bar() executes ret, will jump to exploit code

Exploits Based on Buffer Overflows

Buffer overflow bugs allow remote machines to execute arbitrary code on victim machines

Internet worm

Early versions of the finger server (fingerd) used **gets()** to read the argument sent by the client:

Worm attacked fingerd server by sending phony argument:

finger “exploit-code padding new-return-address”

exploit code: executed a root shell on the victim machine with a direct TCP connection to the attacker.

Exploits Based on Buffer Overflows

Buffer overflow bugs allow remote machines to execute arbitrary code on victim machines

IM War

AOL exploited existing buffer overflow bug in AIM clients
exploit code: returned 4-byte signature (the bytes at some location in the AIM client) to server.

When Microsoft changed code to match signature, AOL changed signature location.

Code Red Exploit Code

Starts 100 threads running

Spread self

Generate random IP addresses & send attack string
Between 1st & 19th of month


Attack www.whitehouse.gov

Send 98,304 packets; sleep for 4-1/2 hours; repeat
Denial of service attack

Between 21st & 27th of month

Deface server's home page

After waiting 2 hours



Avoiding Overflow Vulnerability

```
/* Echo Line */  
void echo()  
{  
    char buf[4]; /* Way too small! */  
    gets(buf);  
    puts(buf);  
}
```

Use library routines that limit string lengths

fgets instead of **gets**

strncpy instead of **strcpy**

Don't use **scanf** with **%s** conversion specification

Use **fgets** to read the string

Or use **%ns** where n is a suitable integer

System-Level Protections

Randomized stack offsets

At start of program, allocate
random amount of space on stack
Makes it difficult for hacker to
predict beginning of inserted code

Nonexecutable code segments

In traditional x86, can mark region
of memory as either “read-only” or
“writeable”

Can execute anything readable

X86-64 added explicit “execute”
permission

```
unix> gdb bufdemo
(gdb) break echo

(gdb) run
(gdb) print /x $ebp
$1 = 0xfffffc638

(gdb) run
(gdb) print /x $ebp
$2 = 0xffffbb08

(gdb) run
(gdb) print /x $ebp
$3 = 0xffffc6a8
```

Stack Canaries

Idea

Place special value (“canary”) on stack just beyond buffer
Check for corruption before exiting function

GCC Implementation

-fstack-protector
-fstack-protector-all

```
unix>./bufdemo-protected  
Type a string:1234  
1234
```

```
unix>./bufdemo-protected  
Type a string:12345  
*** stack smashing detected ***
```

Protected Buffer Disassembly

echo

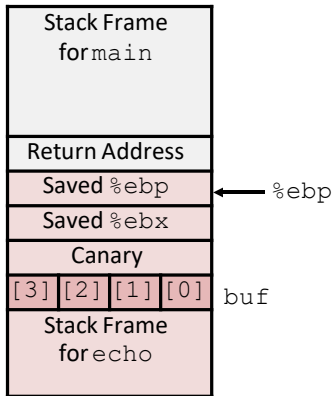
```

804864d: 55                push    %ebp
804864e: 89 e5            mov     %esp,%ebp
8048650: 53              push    %ebx
8048651: 83 ec 14        sub     $0x14,%esp
8048654: 65 a1 14 00 00 00 mov     %gs:0x14,%eax
804865a: 89 45 f8        mov     %eax,0xffffffff8(%ebp)
804865d: 31 c0           xor     %eax,%eax
804865f: 8d 5d f4        lea     0xffffffff4(%ebp),%ebx
8048662: 89 1c 24        mov     %ebx,(%esp)
8048665: e8 77 ff ff ff  call    80485e1 <gets>
804866a: 89 1c 24        mov     %ebx,(%esp)
804866d: e8 ca fd ff ff  call    804843c <puts@plt>
8048672: 8b 45 f8        mov     0xffffffff8(%ebp),%eax
8048675: 65 33 05 14 00 00 00 xor     %gs:0x14,%eax
804867c: 74 05          je      8048683 <echo+0x36>
804867e: e8 a9 fd ff ff  call    804842c <FAIL>
8048683: 83 c4 14        add     $0x14,%esp
8048686: 5b             pop     %ebx
8048687: 5d             pop     %ebp
8048688: c3             ret

```


Setting Up Canary

Before call to gets




```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

```
echo:
...
movl %gs:20, %eax # Get canary
movl %eax, -8(%ebp) # Put on stack
xorl %eax, %eax # Erase canary
...
```

Checking Canary

Before call to gets

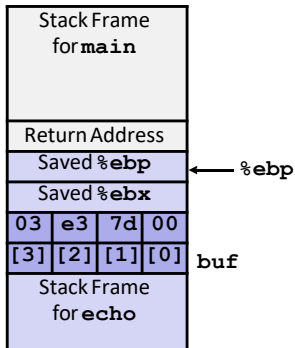


```
/* Echo Line */
void echo()
{
    char buf[4]; /* Way too small! */
    gets(buf);
    puts(buf);
}
```

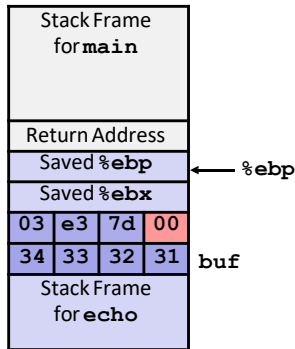
```
echo:
...
movl -8(%ebp), %eax # Retrieve from stack
xorl %gs:20, %eax # Compare with Canary
je .L24 # Same: skip ahead
call stack_chk_fail # ERROR
.L24:
...
```

Canary Example

Before call to gets



Input 1234



```
(gdb) break echo
(gdb) run
(gdb) stepi 3
(gdb) print /x *((unsigned *) $ebp - 2)
$1 = 0x3e37d00
```

Benign corruption!
(allows programmers to make
silent off-by-one errors)

Outline

- 1 Background
- 2 Buffer Overflow
- 3 Integer and Heap Overflow
- 4 Summary

What are the common features of integer overflow vulnerabilities?

an untrusted source

```
unsigned int x = read_int();  
if ( x > 0x7fffffff )  
    abort();  
unsigned int s = x*sizeof(int);  
char* p=malloc(s);  
read_int_into_buf(p, x);
```

an incomplete
check

an integer overflow

a heap overflow
followed

a sensitive
operation

CVE-2008-5238(Xine)

```
.....  
if (version == 4) {  
    const uint16_t sps = _X_BE_16 (this->header+44) ? : 1;  
    this->w          = _X_BE_16 (this->header+42);  
    this->h          = _X_BE_16 (this->header+40);  
    this->cfs        = _X_BE_32 (this->header+24);  
    this->frame_len = this->w * this->h;  
    this->frame_size = this->frame_len * sps;  
    this->frame_buffer = calloc(this->frame_size, 1);  
.....
```

an untrusted source

a sensitive
operation

53 a sensitive operation

CVE-2008-2430(VLC)

```
.....
if( ChunkFind( p_demux, "fmt ", &i_size ) )
{
    msg_Err( p_demux, "cannot find 'fmt ' chunk" );
    goto error;
}
if( i_size < sizeof( WAVEFORMATEX ) - 2 )
{
    msg_Err( p_demux, "invalid 'fmt ' chunk" );
    goto error;
}
stream_Read( p_demux->s, NULL, 8 ); /* Can

/* load waveformatex */
p_wf_ext = malloc( EVEN( i_size ) + 2 );
.....
```

an untrusted source

an incomplete
check

an integer overflow

a sensitive
operation

What's the essential feature of integer overflow vulnerabilities?

an untrusted source

```
unsigned int x = read_int();  
if ( x > 0x7fffffff )  
    abort();  
    unsigned int s = x*sizeof(int);  
char* p=malloc(s);  
read_int_into_buf(p, x);
```

an incomplete
check

an integer overflow

a sensitive
operation

Outline

- 1 Background
- 2 Buffer Overflow
- 3 Integer and Heap Overflow
- 4 Summary

Summary

Software vulnerabilities

- 1 Cannot be avoided (software complexity)
- 2 Memory bugs (buffer overflow, integer overflow) are dangerous
- 3 Tons of research has been carried out to stop memory bugs
- 4 Stack, and integer and heap overflow can be stopped.

References

http://en.wikipedia.org/wiki/Buffer_overflow Smashing
the Stack for Fun and Profit by Aleph One

A Comparison of Buffer Overflow Prevention
Implementations and Weaknesses (Blackhat 2004)

[http://www.cs.cmu.edu/afs/cs/academic/class/15213-
f10/www/lectures/08-machine-advanced.pdf](http://www.cs.cmu.edu/afs/cs/academic/class/15213-f10/www/lectures/08-machine-advanced.pdf)

...